

THE Q COLLECTIVE LLC

d/b/a Doctrine Policy Group · Lakewood, Colorado · theqcollective.org · team@theqcollective.org

October 2, 2026

Attorney General Philip J. Weiser

Colorado Department of Law

Submitted via the ADMT and Chatbot Safety Act rulemaking comment portal, coag.gov/ai

Re: Comments on Proposed Rules Implementing the Automated Decision-Making Technology in Consequential Decisions Act (SB 26-189) and the Conversational Artificial Intelligence Services Act (HB 26-1263), 4 CCR 904-6

Dear Attorney General Weiser and Staff of the Colorado Department of Law:

The Q Collective LLC, operating as Doctrine Policy Group, respectfully submits these comments on the proposed rules published August 11, 2026.

We are a Colorado civic policy organization based in Lakewood. Our work is returning governmental power to constituents. We have no client in this proceeding and no commercial interest in its outcome. We file on behalf of the people of Colorado — the renters, applicants, borrowers, patients, students, and employees who will one day receive an adverse outcome and be told that an automated system helped produce it.

We recognize the Department's genuine openness in this process, and the seriousness of the comments already filed. We write because the record before you is weighted heavily toward the parties who will be regulated, and thinly toward the parties the statute was written to protect.

SUMMARY OF RECOMMENDATIONS

- bullet Adopt Material Influence Standard 1. Standard 2 would raise the threshold the General Assembly set. (§ I)
- bullet Retain Rule 6.6(A)(3) as drafted. Adverse-outcome notices must name data sources. Category-level disclosure makes the statutory correction right unusable. (§ II)
- bullet Retain Rules 7.7(C)(1), 7.7(C)(5), and 7.7(D)(3). Reviewer independence, no automated assistance in reviewing automated output, and the Department's ability to draw inferences from affirmance patterns. (§ III)
- bullet Retain Rule 6.2 (two-channel delivery) and Rule 6.6(B)(1) (ten business days). (§ IV)
- bullet Decline to adopt a post-finalization grace period, and decline to deem ECOA/FCRA compliance sufficient for the Act's adverse-outcome disclosures. (§ IV)

I. THE DEPARTMENT SHOULD ADOPT MATERIAL INFLUENCE STANDARD 1

A. Standard 2 is not a clarification of the statute. It is a revision of it.

The Act defines "materially influence" to mean that the ADMT output is a non-*de minimis* factor affecting the outcome of a consequential decision — including by constraining, ranking, scoring,

recommending, classifying, or otherwise meaningfully altering how that decision is made. C.R.S. § 6-1-1701(13).

Standard 1 gives that language its plain effect. A factor is *de minimis* only where its impact is trifling, trivial, or incidental, and a factor does not become *de minimis* merely because other factors carried more weight.

Standard 2 replaces the statutory phrase with a substantial-factor test, and would permit an output to be treated as *de minimis* whenever independent factors played a significantly larger role.

The General Assembly had the vocabulary of substantiality available to it. Colorado law uses substantial-factor language in other contexts, and the legislature did not choose it here. It chose the *de minimis* framing, which is a floor, not a weighing test. Reading a floor as a weighing test is not clarification within the Department's rulemaking authority. It is amendment of the trigger, accomplished through regulation rather than legislation.

B. What each standard produces in practice.

Consider an applicant for an apartment in Lakewood. A tenant-screening product ingests her rental history, court records, and credit data, and returns a risk tier placing her in the bottom quartile. The landlord declines her, and can also point to a gap in her rental history that he says he weighed more heavily.

Under Standard 1, the output constrained and classified her. The Act applies. She receives notice, learns what data drove the tier, and can correct a record if it is wrong.

Under Standard 2, the landlord's other factor arguably played a significantly larger role. The Act does not apply. She receives nothing. She does not learn she was scored, does not learn what the score rested on, and has no route to correct an error that will follow her to the next application and the one after that.

The underlying conduct is identical in both cases. The difference is solely whether Colorado law reaches it.

C. Standard 1 does not create liability for genuinely independent decisions.

We have read the objection that Standard 1 conflates correlation with causation — that it captures cases where a human independently reached the same conclusion the system did.

Standard 1 already answers this, and generously. Its presumption is rebuttable, and a deployer may rebut it by showing any one of the following: the decision-maker recorded an independent judgment before seeing the output; the decision-maker did not view the output at all; the decision reached was inconsistent with the output; or the decision rested on previously documented and consistently applied dispositive factors.

Those are four separate exits, and each is within the deployer's own control and recordkeeping. A deployer who genuinely decided independently is able to document that it did. A deployer who cannot

produce any such documentation has not shown independent judgment — it has asserted it.

D. The burden belongs with the party that holds the record.

This is the difference that matters most between the two standards, and it is not a difference of degree.

Standard 1 asks the deployer a question about its own conduct: what did your decision-maker do, and can you show it. The deployer chose the tool, negotiated the contract, configured the workflow, and holds the file.

Standard 2 shifts the inquiry to how much other information existed and how heavily it weighed — a question the consumer cannot investigate, cannot access, and cannot contest. It converts the threshold into a matter the regulated party alone can characterize.

Between a deployer who must keep a record and a consumer who must prove a negative about a file she has never seen, a remedial statute points one direction. We ask the Department to follow it.

II. RULE 6.6(A)(3) SHOULD BE RETAINED: DATA SOURCES MUST BE NAMED

A. The proposal before you.

Rule 6.6(A)(3) requires a deployer to identify, by name, each source of the data used in a decision — including specific data brokers, databases, courts, social media companies, schools, and employers.

Commenters have asked the Department to replace names with categories: "consumer reporting agencies," "public records," "application data." The stated grounds are trade-secret protection, competitive harm, and the difficulty deployers face in identifying upstream sources.

We ask the Department to reject that substitution.

B. The federal record already establishes that these errors are common, consequential, and correctable only against a named company.

The Department does not have to take the frequency of these errors on faith. Two federal agencies have measured it.

The errors are common. The Federal Trade Commission's congressionally mandated study of credit report accuracy found that one in five consumers had an error on at least one of their three credit reports, and one in four identified errors that might affect their credit scores. Five percent had errors serious enough to result in less favorable terms.

Consumers who challenge them are usually right. In that same study, four out of five consumers who filed disputes obtained some modification to their report. That figure deserves the Department's attention when weighing arguments that a broader disclosure duty would generate burdensome or meritless disputes. The federal evidence points the other way: when consumers can locate an error and identify who holds it, their challenges succeed far more often than not.

The tenant screening context is worse, not better. The Consumer Financial Protection Bureau's analysis of consumer complaints regarding tenant background checks identified more than 16,000 complaints about incorrect information, and approximately 4,500 complaints specifically about obstacles consumers encountered in correcting errors. The CFPB documented negative information belonging to a different person, outdated records that were never removed, and inaccurate arrest, criminal, and eviction entries that persisted after dispute.

And the cost of each denial is borne up front. The CFPB reports that 68% of renters pay an application fee. A Coloradan rejected on the basis of an erroneous record does not simply lose the unit. She pays again at the next application, and the next, carrying the same uncorrected record with her, because nothing in the process tells her it exists.

C. Category-level disclosure is not a lighter version of this right. It is the absence of one.

The Act grants Coloradans a right to correct inaccurate personal data used against them. C.R.S. § 6-1-1705.

That right operates against an identified entity. It cannot operate against a category. There is no shared dispute portal for "consumer reporting agencies." There is no correction letter that reaches "public records." Each company maintains its own dispute channel, its own identity verification requirements, and its own record — and no consumer who knows only the category can reach any of them.

Return to the FTC's four-in-five figure. Not one of those consumers obtained a correction by writing to a category. Every one of them succeeded because a named company was identifiable and reachable. That is the mechanism the proposed substitution would remove.

A consumer told that "a consumer reporting agency" supplied the record that cost her the apartment has been given information she cannot act on. She has been informed and left without recourse in the same sentence.

We would also note the Department is not without concrete illustration. In *Louis v. SafeRent Solutions*, rental applicants challenged the use of a proprietary tenant screening score and reached a class settlement in 2024. The relevant point for this rulemaking is not the merits of that case but how the applicants learned an algorithmic score had been applied to them at all: through federal litigation and discovery, not through any notice they received at the time of denial. Where litigation is the only reliable route to learning that an automated system evaluated you, the disclosure threshold and the disclosure content must both be set so that ordinary applicants are not dependent on it.

This is the load-bearing connection in the proposed rules. Strip the name out of Rule 6.6(A)(3) and § 6-1-1705 does not weaken. It stops functioning. The right remains on the page with no door to knock on.

D. Category-level disclosure would also sever Coloradans from remedies that exist outside this Act.

The consequence reaches further than the correction right in § 6-1-1705, and we ask the Department to weigh this carefully.

Many of the companies supplying the data behind consequential decisions in Colorado are consumer reporting agencies. Consumers already hold rights against them under the federal Fair Credit Reporting Act, including a private right of action for inaccurate reporting. Other remedies exist under federal and state fair housing and fair lending law. None of those depend on the ADMT Act, and none of them are before the Department in this rulemaking.

But every one of them requires an identifiable party.

A consumer cannot send an FCRA dispute to a category. She cannot name "a consumer reporting agency" as a defendant. She cannot direct a fair housing complaint at "public records." The rights survive the Chamber's proposal in form and become unexercisable in fact, because the only route to any of them runs through knowing who to name.

We do not suggest the Department is charged with enforcing those other laws. We suggest the opposite: precisely because those remedies lie outside the Department's authority, the Department should be reluctant to adopt a disclosure standard that quietly forecloses them. A rule that makes an actor unnameable does not merely limit this statute. It limits every statute that would otherwise reach that actor.

E. The identity of a vendor is not a trade secret.

The proprietary assets in an automated decision system are the model, the weighting, the training data, and the algorithm. Rule 6.6(A)(3) reaches none of them.

Disclosing that a landlord purchased a screening report from a specific named company reveals nothing about how that company scores, what it weighs, or how it was built. It reveals a commercial relationship. Commercial relationships are routinely disclosed under existing law: the Fair Credit Reporting Act has required named-source adverse action notices for decades, and the consumer reporting industry has not been harmed by it.

F. A deployer that cannot name its sources has described a procurement failure, not a disclosure burden.

Where a deployer using automated tools to decide Coloradans' access to housing, credit, or employment cannot say where the underlying data came from, the appropriate regulatory response is not to excuse the disclosure. It is to require the deployer to obtain that information from its vendors.

That obligation flows naturally through the contracting chain, and it is precisely what the Department's own questions regarding multiparty arrangements and ADMT vendors contemplate. We encourage the Department to close that loop expressly: a deployer remains responsible for source identification regardless of vendor involvement, and vendors should be obligated to supply what the deployer needs to comply.

III. THE MEANINGFUL HUMAN REVIEW SAFEGUARDS IN RULE 7.7 SHOULD BE RETAINED

Human review is the consumer's only mechanism for obtaining a second look at an automated outcome. Three proposed deletions would leave that review intact in name and unverifiable in substance.

A. Rule 7.7(C)(1) — reviewer independence.

It has been proposed that the Department remove the requirements that the reviewer not be a subordinate of the original decision-maker and be shielded from retaliation.

Review conducted by the original decision-maker's subordinate is not review. It is ratification with additional paperwork. A reviewer who must work under that supervisor the following morning does not hold the practical authority to override that the statute presupposes, whatever the org chart says.

We take seriously the concern that this is difficult for very small employers. The answer to that is a narrow accommodation scaled to headcount — not the removal of independence as a requirement for everyone, including deployers processing tens of thousands of decisions a year.

B. Rule 7.7(C)(5) — no automated assistance in the review.

It has been proposed that "meaningful" should not be read to mean "exclusively human."

But the review exists because an automated output is being contested. Permitting a system of the same class to assist in evaluating that output returns the consumer to the machine she appealed from, one step removed. Whatever else meaningful human review requires, it cannot mean the algorithm gets a second vote on its own work.

C. Rule 7.7(D)(3) — inferences from affirmance.

It has been proposed that the Department state affirmatively that reaching the same conclusion as the ADMT does not indicate review was not meaningful.

We do not object to the Department clarifying that a single affirmance is not itself proof of failure. We object to language that forecloses the inference entirely.

If a deployer's human reviewers affirm the automated output in substantially every case, that pattern is evidence, and the Department should preserve its ability to treat it as such. We suggest the Department instead require deployers to retain review outcome rates under Rule 7.7(H), and confirm that a near-total affirmance rate may support a finding that review was not meaningful.

D. Rule 7.7(H) — record retention.

It has been argued that detailed review records create litigation exposure.

Records that would prove damaging in litigation are records of conduct that was damaging to a consumer. The statute already sets three-year retention, and the legislature already classified violations as deceptive trade practices. That exposure is not a defect in the rule. It is the enforcement mechanism operating as designed.

We would add a broader point that applies to Rule 6.6(A)(3) and Rule 7.7(H) together. The General Assembly placed enforcement of this Act with the Attorney General and did not create a private right of action. That makes the Department's records and disclosure rules the only mechanism by which the facts of an automated decision are preserved at all. Whether those facts are ever used — by this Department, by a federal regulator, by a consumer exercising rights under another statute, or by a future General Assembly assessing how this framework performed — depends entirely on whether the record exists. A disclosure rule is, in practice, an evidence rule. We ask the Department to keep that in view when weighing requests to reduce what must be recorded and what must be disclosed.

IV. TIMING AND ACCESS PROVISIONS

A. Rule 6.2 — two-channel delivery. Retain.

Single-channel notice fails precisely the population this Act protects. People who have recently moved, lost access to an email account, or are between housing situations are the same people most likely to be on the receiving end of an adverse automated outcome. The marginal cost of a second delivery channel is measured in cents. The cost of a notice that never arrives is a statutory right that expires in silence.

B. Rule 6.6(B)(1) — ten business days. Retain.

It has been proposed that supplemental information requests be given thirty days, by analogy to the initial disclosure.

But the consumer making a supplemental request is already inside a closing window. Appeal deadlines, lease deadlines, hiring timelines, and rescreening cycles do not pause while she waits. Thirty days for supplemental information, stacked behind thirty days for the initial notice, can run past sixty days — well past the point at which correcting the record changes any outcome for her.

If operational feasibility is a genuine concern, we would not object to a limited extension mechanism for documented complex cases. We would object to a categorical doubling of the default.

C. Post-finalization grace period. Decline.

Both Acts were signed May 14, 2026, with a January 1, 2027 effective date. Proposed rules issued August 11, 2026 gave nearly five months of advance visibility into the Department's direction.

A grace period running from finalization would move the first day of actual protection for Colorado consumers into 2027 or beyond, for a statute the General Assembly set to take effect on January 1. We respectfully submit that the implementation runway was set by the legislature, and that the Department was not asked to extend it.

D. ECOA and FCRA deemed compliance. Decline as proposed.

Federal adverse action notices were designed around credit reports. They were not designed around automated systems that score, rank, and classify on inputs far broader than credit history, and they do not deliver what this Act requires: identification of the ADMT, the principal reasons, the data relied on,

correction instructions, and the route to human review.

Deeming ECOA and FCRA compliance sufficient would exempt a large share of consequential decisions from the Act's central disclosures by substituting the narrower federal standard for the broader state one the legislature enacted.

If the Department wishes to reduce duplicative paperwork, the narrower path is available: permit a single combined notice that satisfies both regimes. That eliminates the duplication without eliminating the protection.

V. CONCLUSION

The Department has been asked, in detail and at length, to narrow this rule set. Many of those requests are made in good faith by parties with real operational concerns, and some of them merit accommodation.

But we ask the Department to keep one fact in view while weighing them.

The Coloradans this Act was written for are not in this docket. They do not have counsel, do not file comments, and will never read 4 CCR 904-6. They will encounter these rules exactly once — on the day they are denied housing, a job, a loan, or care they needed, and they want to know why.

On that day, every provision addressed above is the difference between an answer they can act on and a notice they can only put in a drawer.

We appreciate the Department's work on this rulemaking and the care evident in the proposed rules. We are glad to provide any additional information that would assist the Department.

Respectfully submitted,

maveriq B. Jackson

Founder

The Q Collective LLC / Doctrine Policy Group

Lakewood, Colorado

Filed by The Q Collective on behalf of the people of Colorado. No client, no retainer, and no compensated interest in this proceeding.

team@theqcollective.org · theqcollective.org

SOURCES CITED

^{bullet}Federal Trade Commission, Report to Congress Under Section 319 of the Fair and Accurate Credit Transactions Act of 2003 (credit report accuracy study) — one in five consumers with an error on at least one report; one in four identifying errors that might affect scores; five percent with errors capable of producing less favorable terms; four in five disputing consumers obtaining a modification.

- bullet Consumer Financial Protection Bureau, reports on tenant background checks — 16,000+ complaints regarding incorrect information; approximately 4,500 regarding obstacles to correcting errors; 68% of renters paying application fees.
- bullet *Louis v. SafeRent Solutions, LLC*, No. 1:22-cv-10800 (D. Mass.), class settlement 2024.
- bullet C.R.S. § 6-1-1701(13); C.R.S. § 6-1-1704; C.R.S. § 6-1-1705.
- bullet Proposed Automated Decision-Making Technology and Conversational Artificial Intelligence Service Rules, 4 CCR 904-6, filed August 11, 2026, and the accompanying Notice of Rulemaking Hearing.